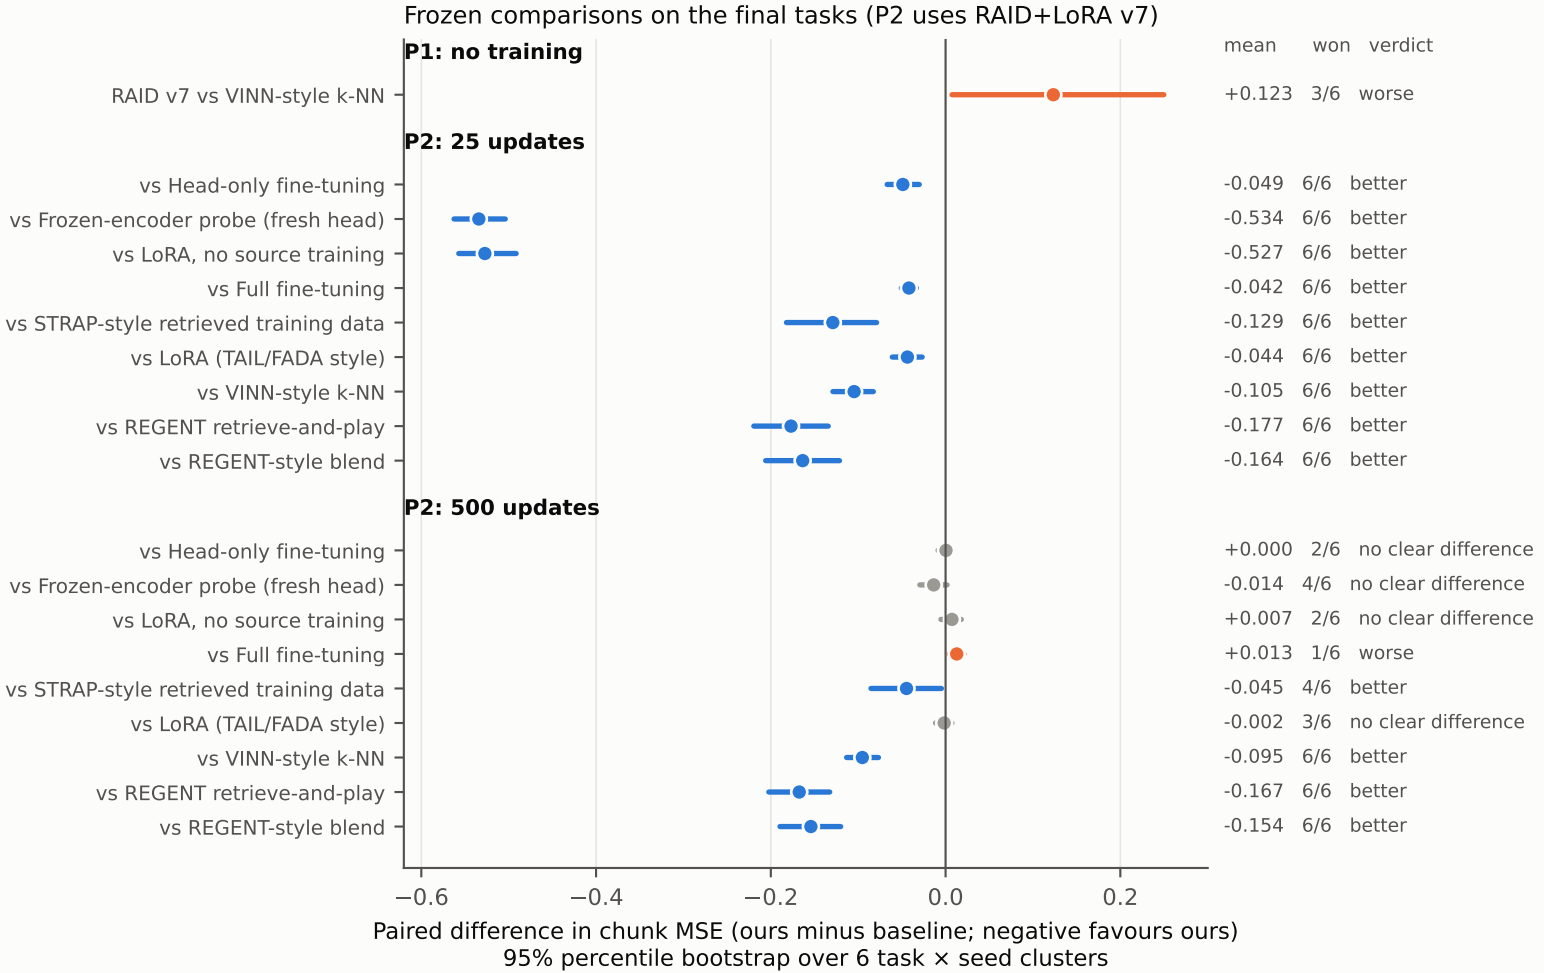
Frozen comparisons on the final tasks (P2 uses RAID+LoRA v7)
P1: no training
RAID v7 vs VINN-style k-NN
P2: 25 updates
vs Head-only fine-tuning
vs Frozen-encoder probe (fresh head)
vs LoRA, no source training
vs Full fine-tuning
vs STRAP-style retrieved training data
vs LoRA (TAIL/FADA style)
vs VINN-style k-NN
vs REGENT retrieve-and-play
vs REGENT-style blend
P2: 500 updates
vs Head-only fine-tuning
vs Frozen-encoder probe (fresh head)
vs LoRA, no source training
vs Full fine-tuning
vs STRAP-style retrieved training data
vs LoRA (TAIL/FADA style)
vs VINN-style k-NN
vs REGENT retrieve-and-play
vs REGENT-style blend
mean won verdict
+0.123 3/6 worse
-0.049 6/6 better
-0.534 6/6 better
-0.527 6/6 better
-0.042 6/6 better
-0.129 6/6 better
-0.044 6/6 better
-0.105 6/6 better
-0.177 6/6 better
-0.164 6/6 better
+0.000 2/6 no clear difference
-0.014 4/6 no clear difference
+0.007 2/6 no clear difference
+0.013 1/6 worse
-0.045 4/6 better
-0.002 3/6 no clear difference
-0.095 6/6 better
-0.167 6/6 better
-0.154 6/6 better
−0.6
−0.4
−0.2
0.0
0.2
Paired difference in chunk MSE (ours minus baseline; negative favours ours)
95% percentile bootstrap over 6 task × seed clusters