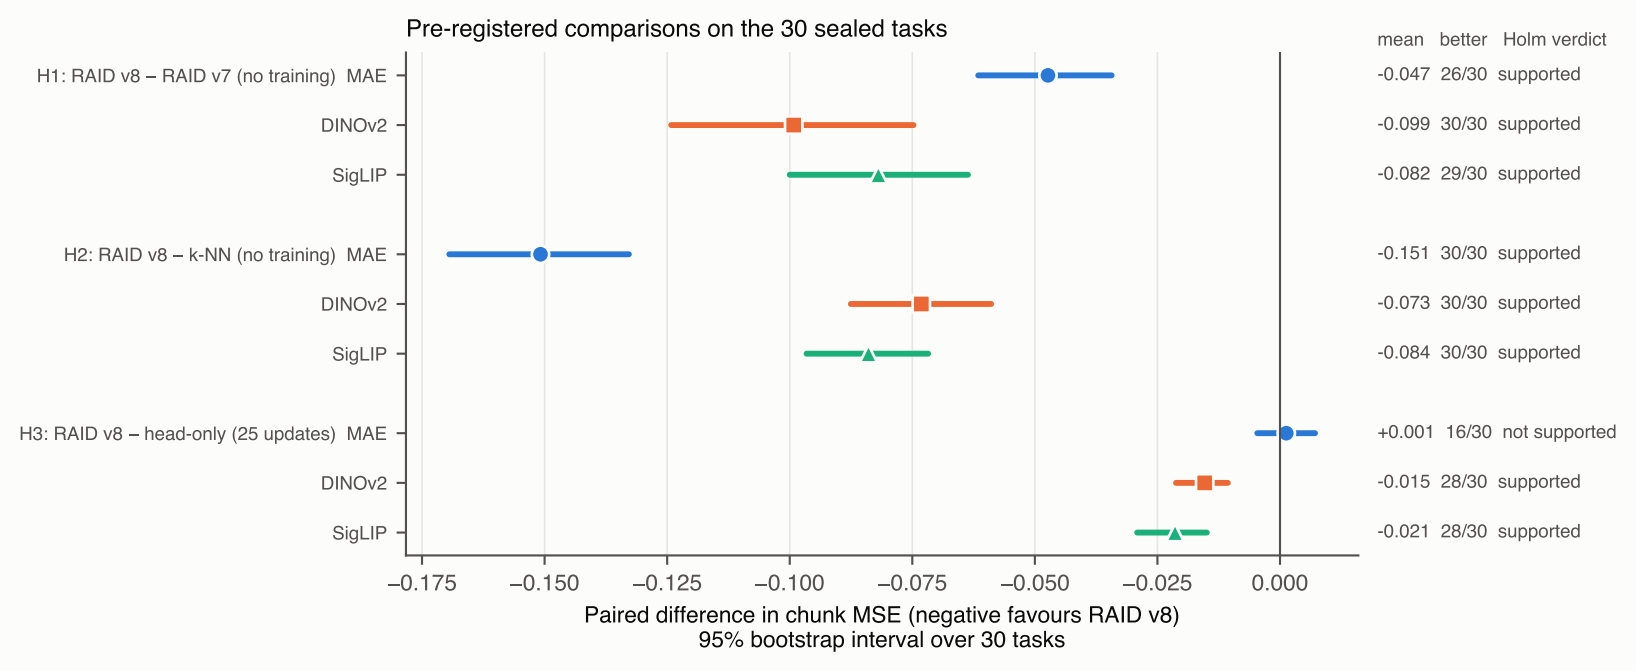
Pre-registered comparisons on the 30 sealed tasks
mean better Holm verdict
H1: RAID v8 – RAID v7 (no training) MAE
-0.047 26/30 supported
DINOv2
-0.099 30/30 supported
SigLIP
-0.082 29/30 supported
H2: RAID v8 – k-NN (no training) MAE
-0.151 30/30 supported
DINOv2
-0.073 30/30 supported
SigLIP
-0.084 30/30 supported
H3: RAID v8 – head-only (25 updates) MAE
+0.001 16/30 not supported
DINOv2
-0.015 28/30 supported
SigLIP
-0.021 28/30 supported
−0.175
−0.150
−0.125
−0.100
−0.075
−0.050
−0.025
0.000
Paired difference in chunk MSE (negative favours RAID v8)
95% bootstrap interval over 30 tasks